How can you ensure there are at least 3 Pod instances (example) are always available and running at point in time?

srinathchalla@outlook.com

Replication Controller

Concept

Udemy

Objectives

Concept

a. Replication Controller Overview
b. Advantages
c. Manifest file

Review Demo

a. Deploy application with Replication Controller
b. Display and validate
c. Scale up and Scale down the application
d. Clean up

srinathchalla@outlook.com

Replication Controller

- Ensures that a specified number of pods are running at any time
    a. If there are excess Pods, they get killed and vice versa
    b. New Pods are launched when they get fail, get deleted or terminated
- Replication Controllers and Pods are associated with “ labels ”
- Creating a “rc” with count of 1 ensure that a pod is always available

srinathchalla@outlook.com

High Availability

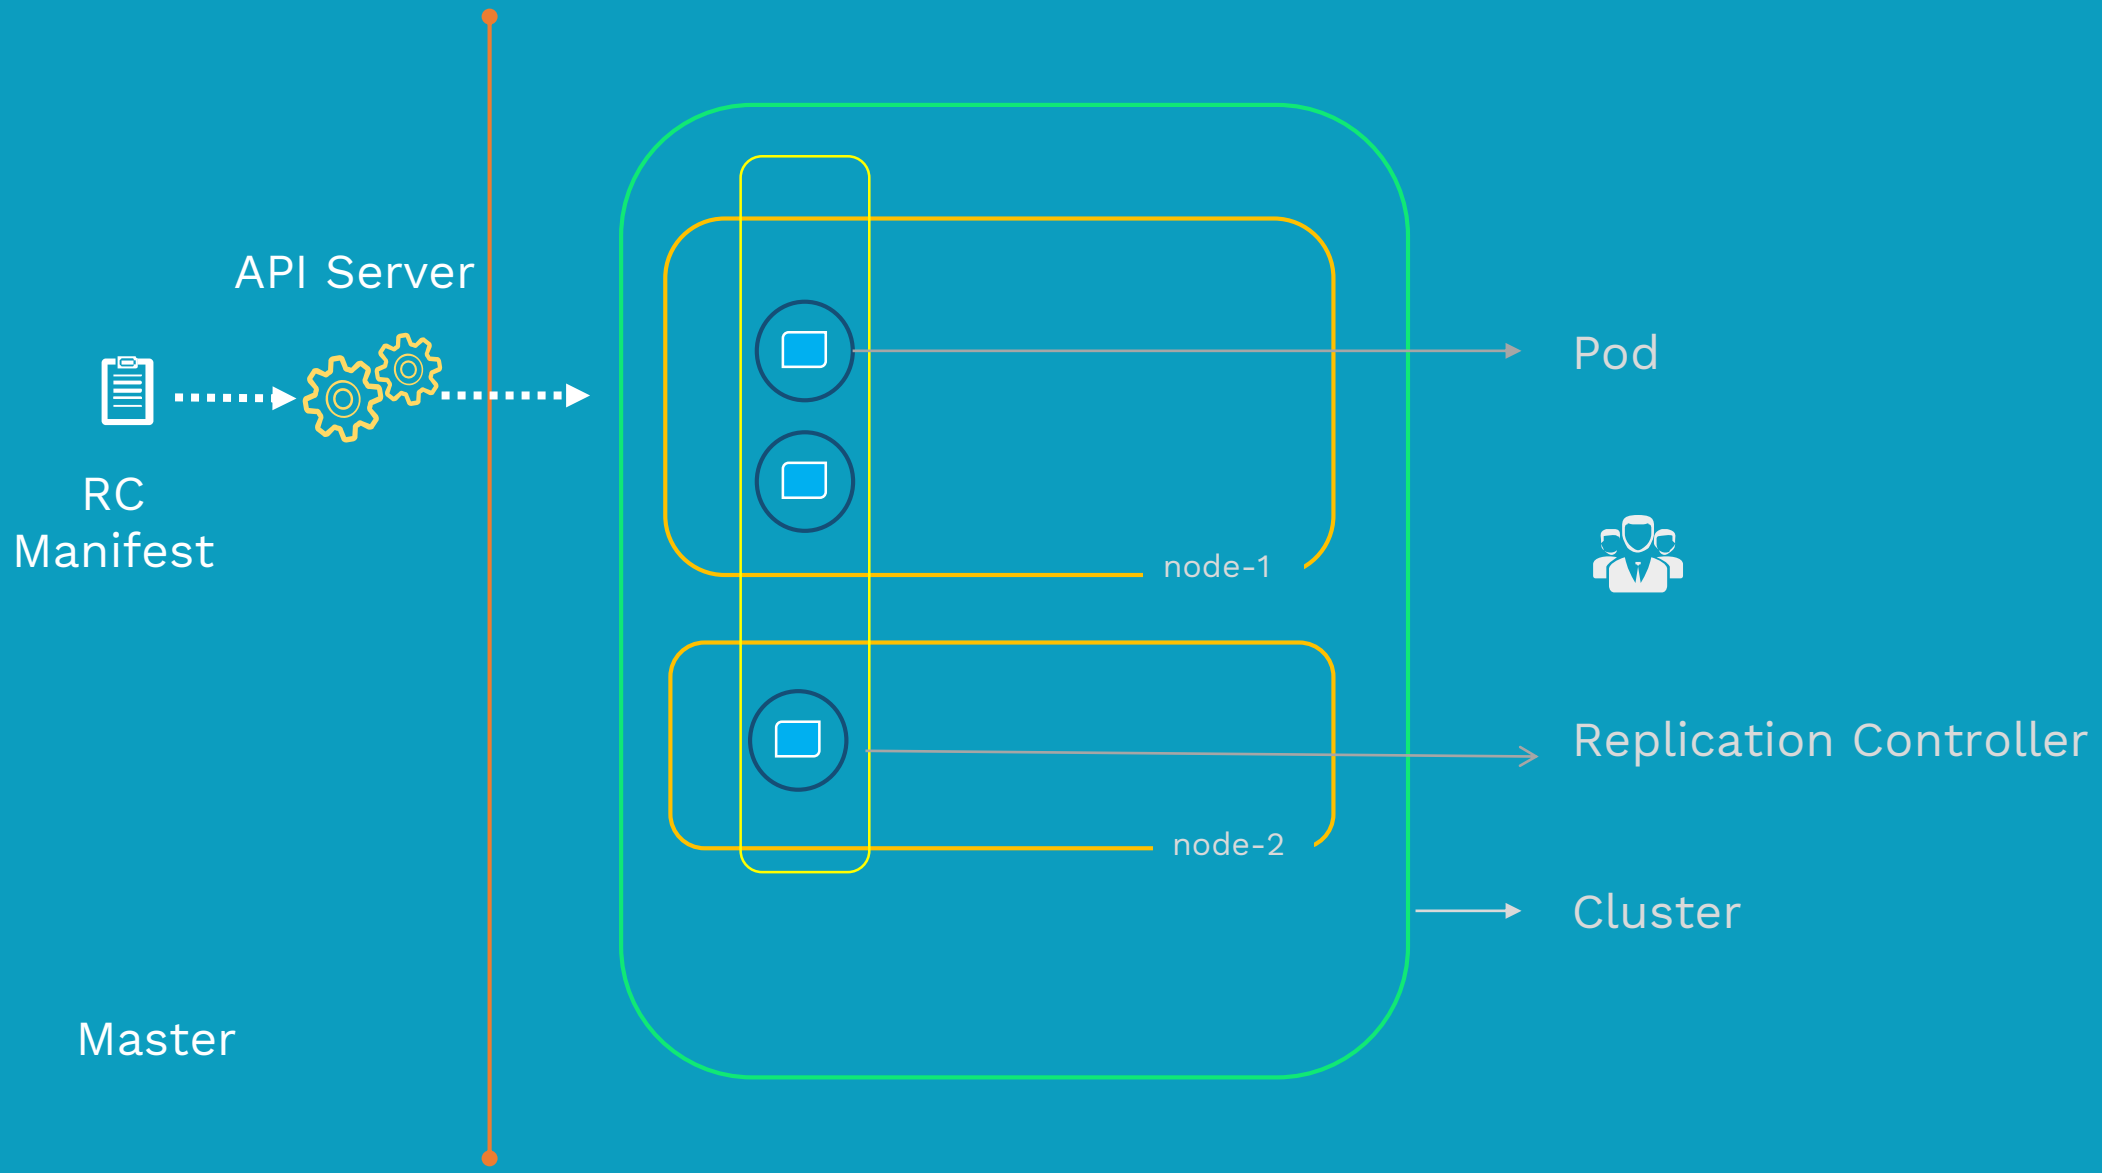

srinathchalla@outlook.com

Load Balancing

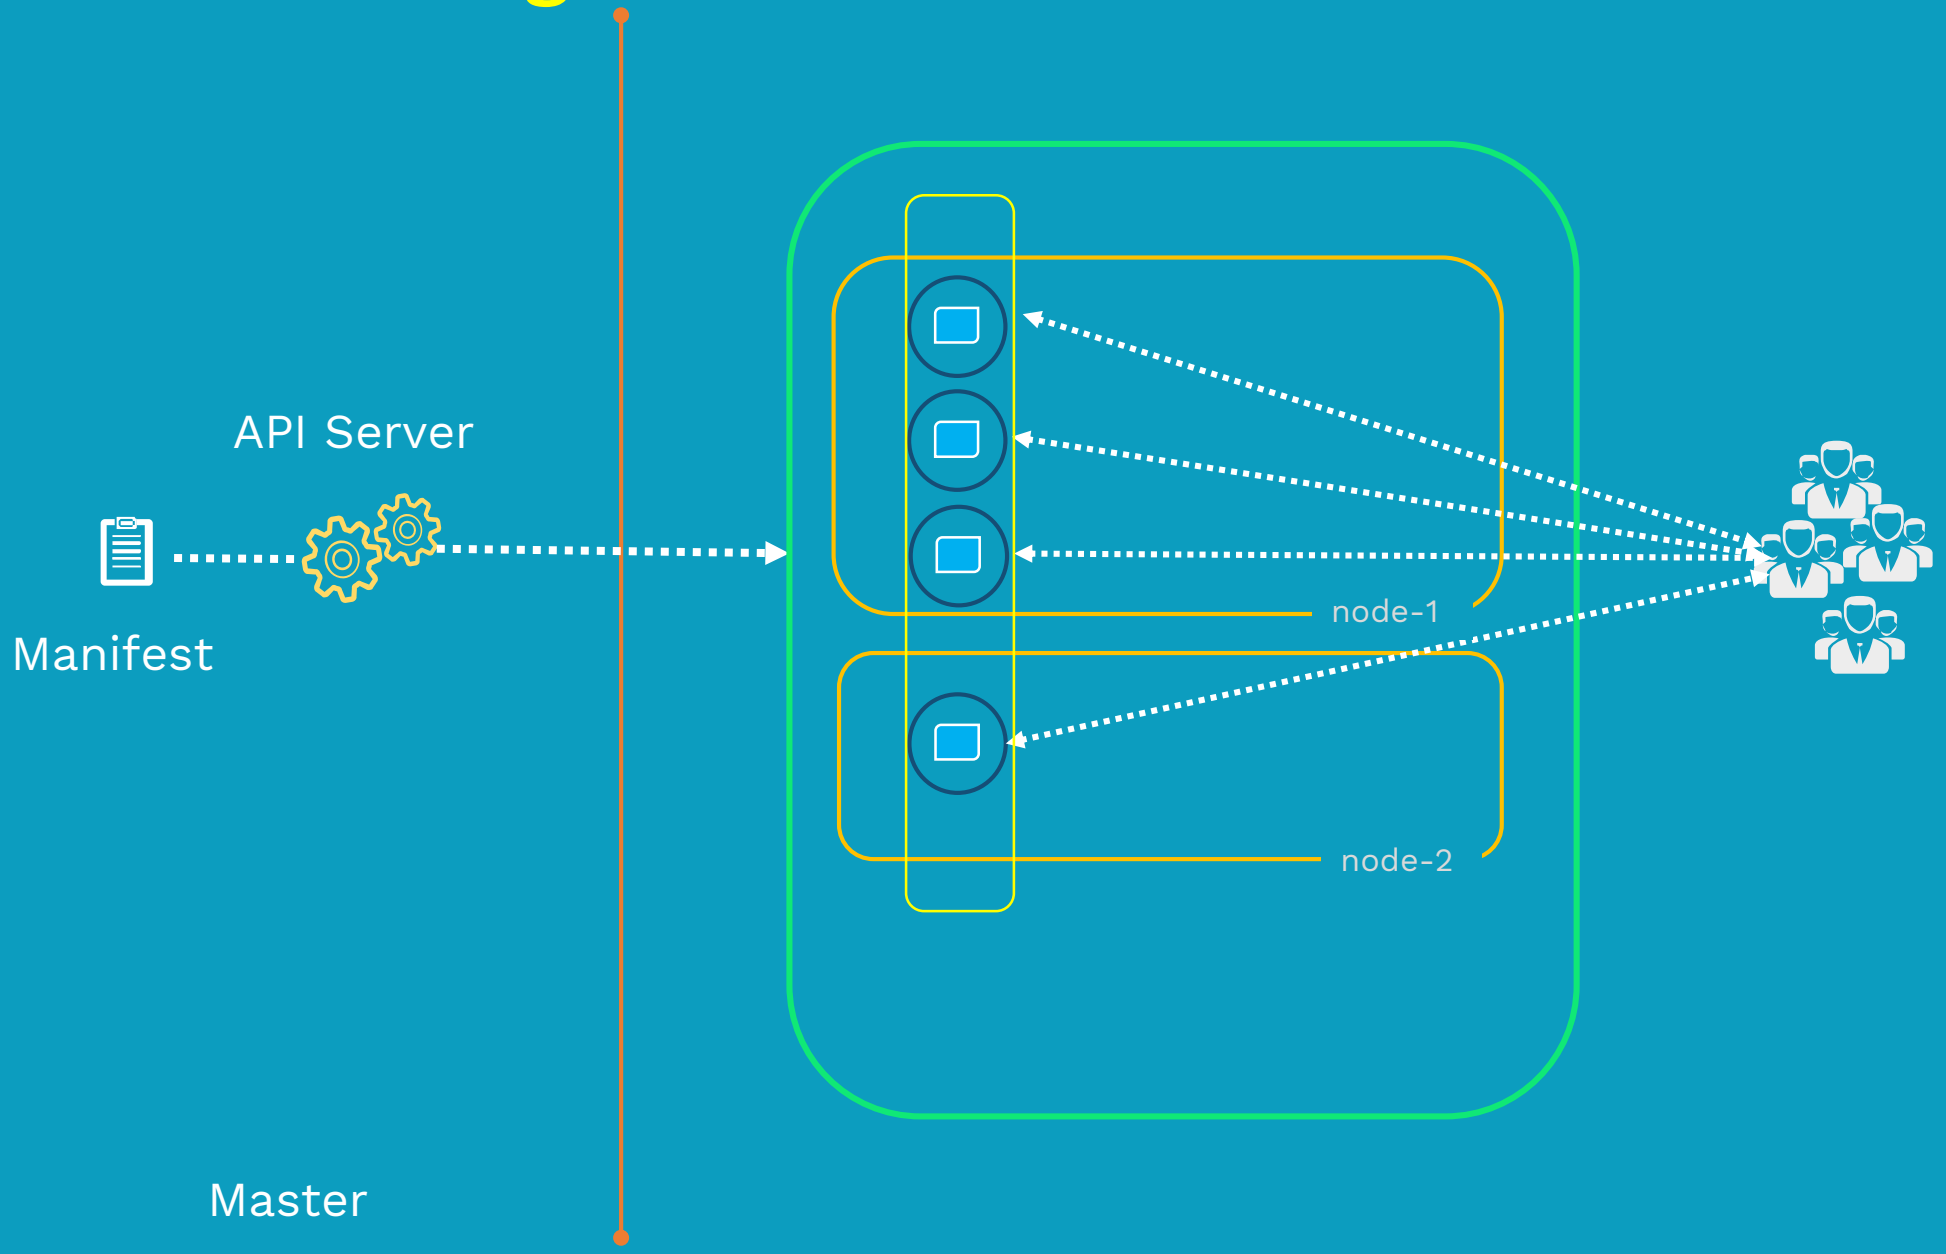

srinathchalla@outlook.com

One last thing

Replication Controller - OLD

ReplicaSet - NEW

Review
Demo

a. Manifest file
b. Deploy app using RC
c. Display and validate RC
d. Test – Node fails
e. Test – Scale up
f. Test – Scale down

srinathchalla@outlook.com

Replication Controller – Config

```yaml
# nginx-rc.yaml
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx-rc
spec:
  replicas: 3
  selector:
    app: nginx-app
  template:
    metadata:
      name: nginx-pod
      labels:
        app: nginx-app
    spec:
      containers:
      - name: nginx-container
        image: nginx
        ports:
        - containerPort: 80
```

Pod def

srinathchalla@outlook.com

Replication Controller – Create & Display

```
[srinath@master ~]$ kubectl create -f nginx-rc.yaml
replicationcontroller/nginx-rc created
```

```
[srinath@master ~]$ kubectl get pods
NAME             READY     STATUS    RESTARTS   AGE
nginx-rc-62vpv   1/1       Running   0          6m
nginx-rc-fk67w   1/1       Running   0          6m
nginx-rc-qk4ph   1/1       Running   0          6m
```

```
[srinath@master ~]$ kubectl get po -l app=nginx-app
NAME             READY     STATUS    RESTARTS   AGE
nginx-rc-62vpv   1/1       Running   0          4m
nginx-rc-fk67w   1/1       Running   0          4m
nginx-rc-qk4ph   1/1       Running   0          4m
```

srinathchalla@outlook.com

Replication Controller – Describe

```
[srinath@master ~]$ kubectl describe rc nginx-rc
Name:         nginx-rc
Namespace:    default
Selector:     app=nginx-app
Labels:       app=nginx-app
Annotations:  <none>
Replicas:     3 current / 3 desired
Pods Status:  3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=nginx-app
  Containers:
   nginx-container:
    Image:        nginx
    Port:         80/TCP
    Host Port:    0/TCP
    Environment:  <none>
    Mounts:       <none>
  Volumes:        <none>
Events:
  Type    Reason            Age   From                    Message
  ----    ------            ----  ----                    -------
  Normal  SuccessfulCreate  25m   replication-controller  Created pod: nginx-rc-62vpv
  Normal  SuccessfulCreate  25m   replication-controller  Created pod: nginx-rc-fk67w
  Normal  SuccessfulCreate  25m   replication-controller  Created pod: nginx-rc-qk4ph
```

srinathchalla@outlook.com

Replication Controller – Node Fail

```
[srinath@master ~]$ kubectl get po -o wide
NAME              READY     STATUS    RESTARTS   AGE       IP             NODE      NOMINATED NODE
nginx-rc-5lb6b    1/1       Running   0          16s       10.240.1.32    node1     <none>
nginx-rc-l64js    1/1       Running   0          16s       10.240.1.33    node1     <none>
nginx-rc-qrqv7    1/1       Running   0          16s       10.240.2.28    node2     <none>
```

```
[srinath@master ~]$ kubectl get nodes
NAME      STATUS     ROLES     AGE       VERSION
master    Ready      master    1d        v1.11.2
node1     Ready      <none>    1d        v1.11.2
node2     NotReady   <none>    1d        v1.11.2
```

```
[srinath@master ~]$ kubectl get po -o wide
NAME              READY     STATUS    RESTARTS   AGE       IP             NODE      NOMINATED NODE
nginx-rc-5lb6b    1/1       Running   0          11m       10.240.1.32    node1     <none>
nginx-rc-l64js    1/1       Running   0          11m       10.240.1.33    node1     <none>
nginx-rc-mtwzs    1/1       Running   0          2m        10.240.1.34    node1     <none>
nginx-rc-qrqv7    1/1       Unknown   0          11m       10.240.2.28    node2     <none>
```

srinathchalla@outlook.com

Replication Controller – Scaling up

```
[srinath@master ~]$ kubectl scale rc nginx-rc --replicas=5
replicationcontroller/nginx-rc scaled
```

```
[srinath@master ~]$ kubectl get rc nginx-rc
NAME       DESIRED   CURRENT   READY     AGE
nginx-rc   5         5         5         9m
```

```
[srinath@master ~]$ kubectl get po -o wide
NAME             READY     STATUS    RESTARTS   AGE       IP              NODE  NOMINATED NODE
nginx-rc-2x2kg   1/1       Running   0          36m       10.240.1.173    node1 <none>
nginx-rc-7kvhl   1/1       Running   0          36m       10.240.2.3      node2 <none>
nginx-rc-g7mwq   1/1       Running   0          33m       10.240.2.42     node2 <none>
nginx-rc-jgt28   1/1       Running   0          33m       10.240.1.3      node1 <none>
nginx-rc-wvmrx   1/1       Running   0          36m       10.240.2.4      node2 <none>
```

srinathchalla@outlook.com

Replication Controller – Scaling down

```
[srinath@master ~]$ kubectl scale rc nginx-rc --replicas=3
replicationcontroller/nginx-rc scaled
```

```
[srinath@master ~]$ kubectl get rc nginx-rc
NAME       DESIRED   CURRENT   READY     AGE
nginx-rc   3         3         3         45m
```

```
[srinath@master ~]$ kubectl get po -o wide
NAME             READY     STATUS    RESTARTS   AGE       IP              NODE  NOMINATED NODE
nginx-rc-2x2kg   1/1       Running   0          36m       10.240.1.173    node1 <none>
nginx-rc-jgt28   1/1       Running   0          33m       10.240.1.3      node1 <none>
nginx-rc-wvmrx   1/1       Running   0          36m       10.240.2.4      node2 <none>
```

srinathchalla@outlook.com

Replication Controller – Delete

```
[srinath@master ~]$ kubectl delete -f nginx-rc.yaml  #kubectl delete rc nginx-rc
replicationcontroller "nginx-rc" deleted
```

```
[srinath@master ~]$ kubectl get rc
No resources found.
```

```
[srinath@master ~]$ kubectl get po -l app=nginx-app
No resources found.
```

srinathchalla@outlook.com

Summary

Concept

a. Ensures that a specified number of pods are running at any time

b. Advantages

Review Demo

a. Manifest file

b. Deploy application with Replication Controller

c. Display and validate

d. Scale up and Scale down the application

e. Clean up

srinathchalla@outlook.com

Coming up...

Demo

Replication Controller

srinathchalla@outlook.com